

Norm AI - Reflective Response Section

What unique challenges do you foresee in developing and integrating AI regulatory agents for legal compliance from a full-stack perspective? How would you address these challenges to make the system robust and user-friendly?

One of the biggest challenges I anticipate in developing and integrating AI regulatory agents is the processing of unstructured data. Both laws and the material to be checked against them will come in a variety of formats, therefore there will be no “one size fits all” solution to processing this data. I can imagine there being a necessity to validate that this data is correctly parsed through both parsing automation, perhaps another AI implementation that gives a parsing confidence rating, and then through human validation when needed.

Another issue I could foresee is the issue when laws overlap between jurisdictions. For example, one requirement might exist at the federal level while another applies at the state level. To handle this, all relevant laws across the jurisdictions would need to be displayed to the user, and inheritance logic should be implemented to show which take precedence. Providing users with the ability to specify the location or context they want their materials checked against would also help filter results according to their needs

There is also the issue of user trust. Legal matters are very sensitive and accuracy is extremely important, so it is necessary the user has full trust in the application. Therefore, it's crucial to minimize hallucinations, (through RAG, as mentioned in the project), deliver confidence ratings (something I know Norm already implements), as well as implement User Feedback loops, so that the user can let the system know when a result is incorrect, and the system can train on that data.